

Quicker Sorting

Dietmar Kühl
dkuhl@bloomberg.net
Real-Time Distribution London
Bloomberg L.P.

Thanks to the sponsors!

Bloomberg



think-cell 



Copyright Notice

© 2016 Bloomberg L.P. Permission is granted to copy, distribute, and display this material, and to make derivative works and commercial use of it. The information in this material is provided "AS IS", without warranty of any kind. Neither Bloomberg nor any employee guarantees the correctness or completeness of such information. Bloomberg, its employees, and its affiliated entities and persons shall not be liable, directly or indirectly, in any way, for any inaccuracies, errors or omissions in such information. Nothing herein should be interpreted as stating the opinions, policies, recommendations, or positions of Bloomberg.

Overview

- quickly sorting out Quick-Sort
- implement and compare variations of Quick-Sort
- turning the concrete implementation generic

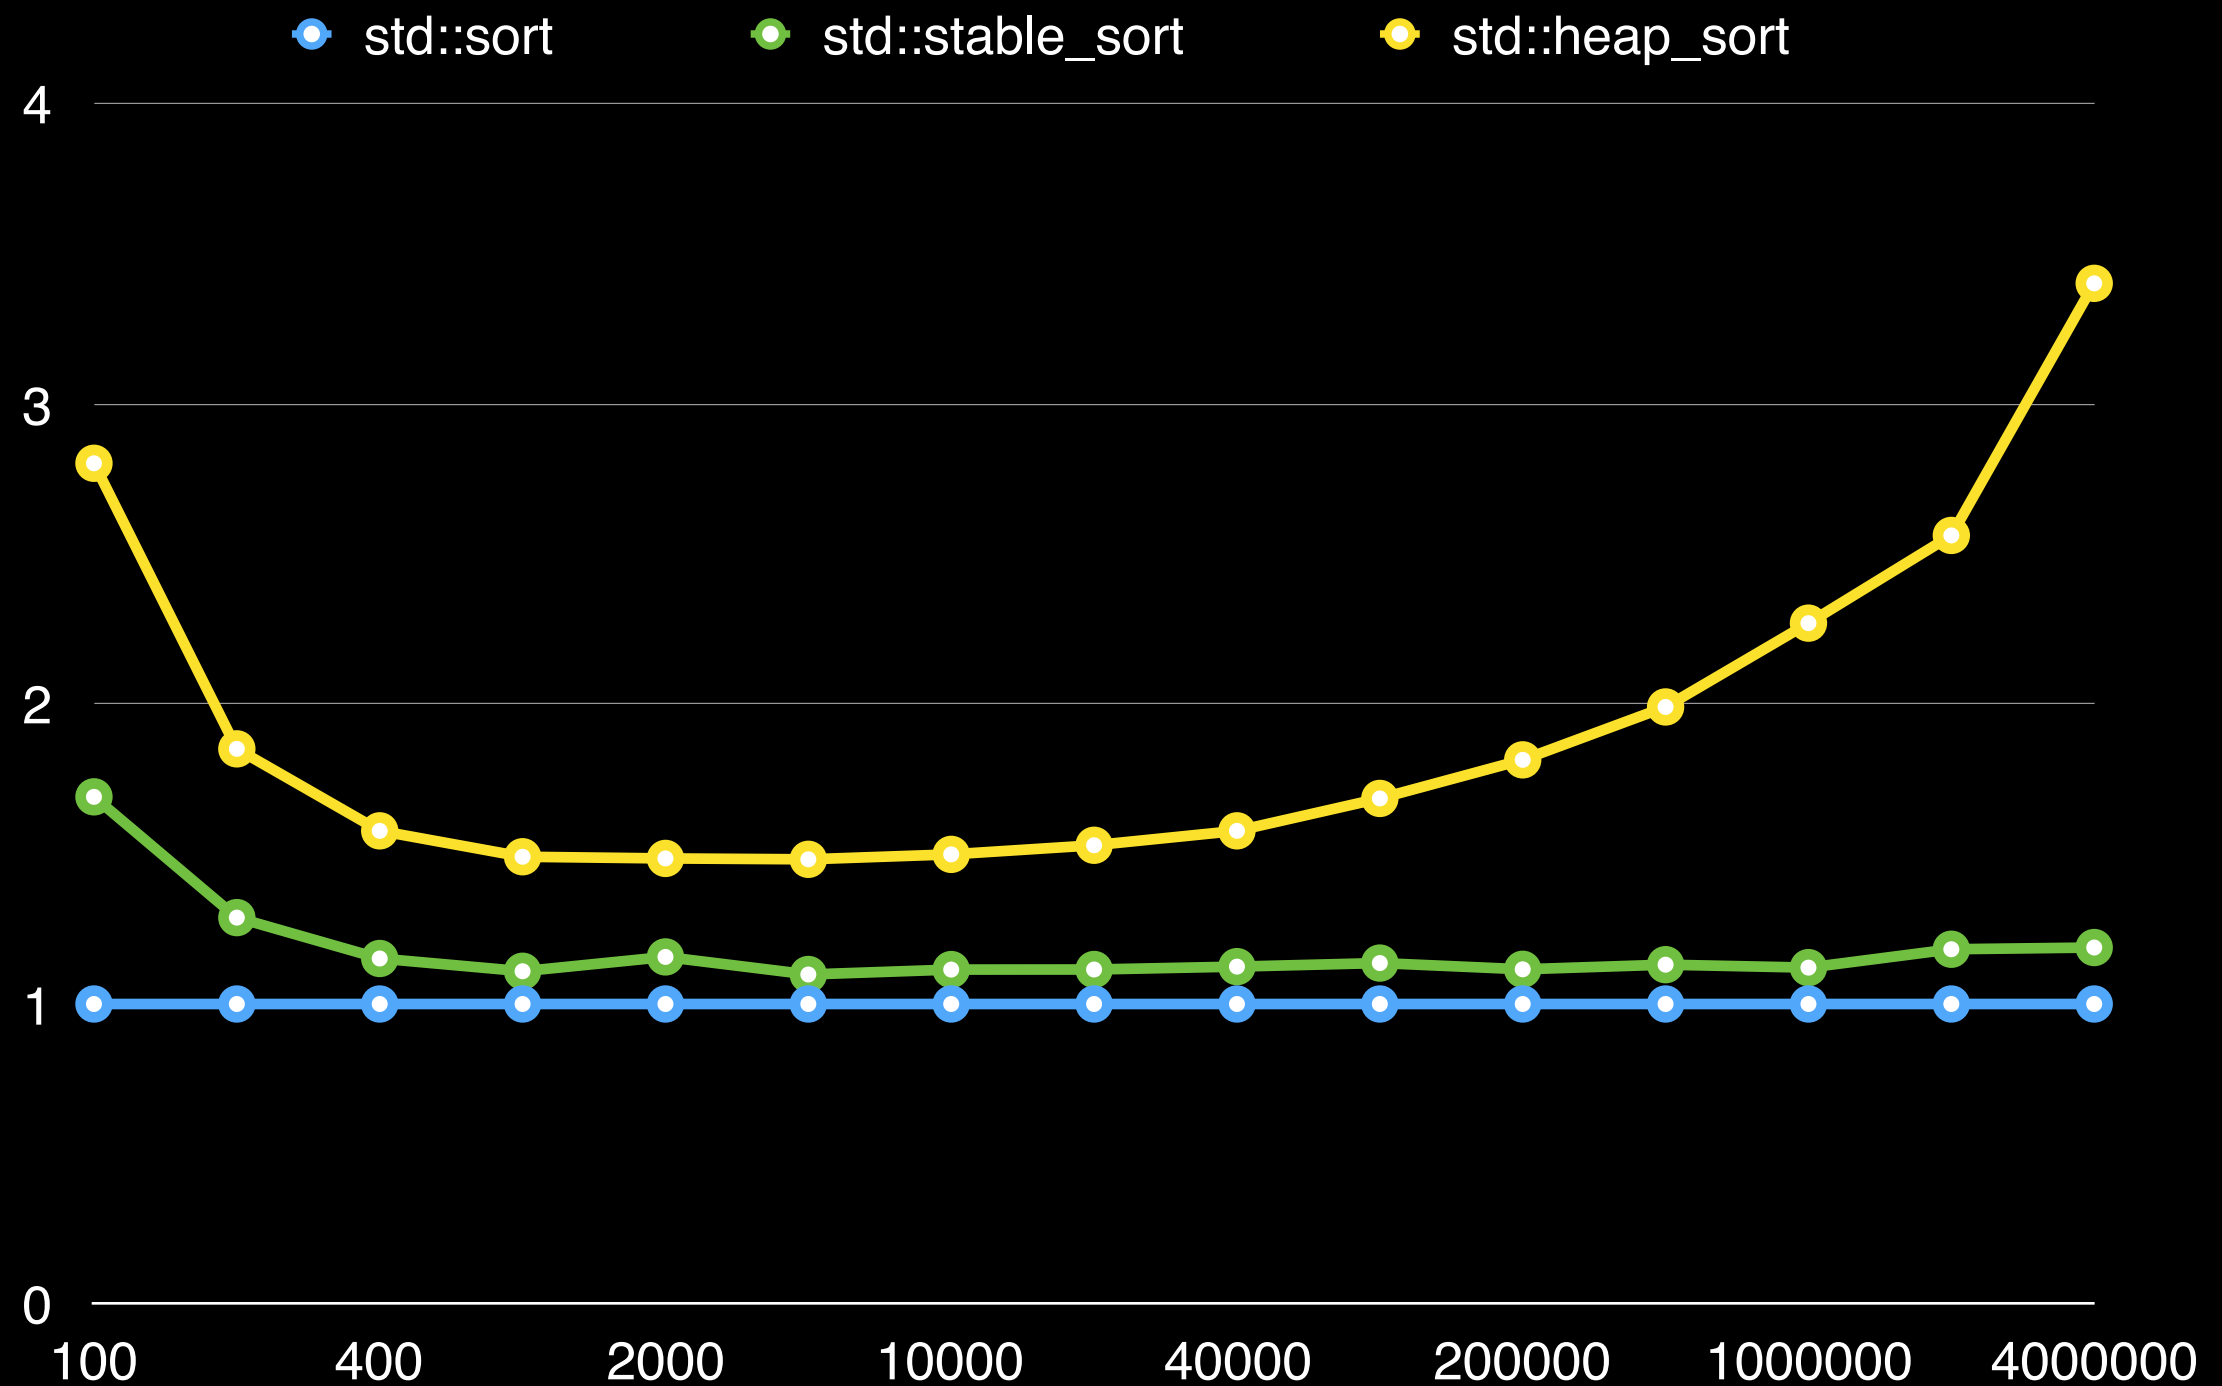
C++ Standard Library

- complexity of C++ Standard Library sort algorithms is $O(n \log n)$
- `std::sort()` - in-place, fast
- `std::stable_sort()` - stable, not in-place
- `std::make_heap()/std::sort_heap()`

Data Set

- multiple random sequence
 - with varying degrees of equal values
- one each ascending/descending sequence
- one sequence with all the same value

Results



Quick-Sort

input



Quick-Sort

input



pivot



Quick-Sort


input



pivot



$x < p$



Quick-Sort


input




pivot



$x < p$



partition



Quick-Sort


input




pivot




$x < p$



partition



recurse

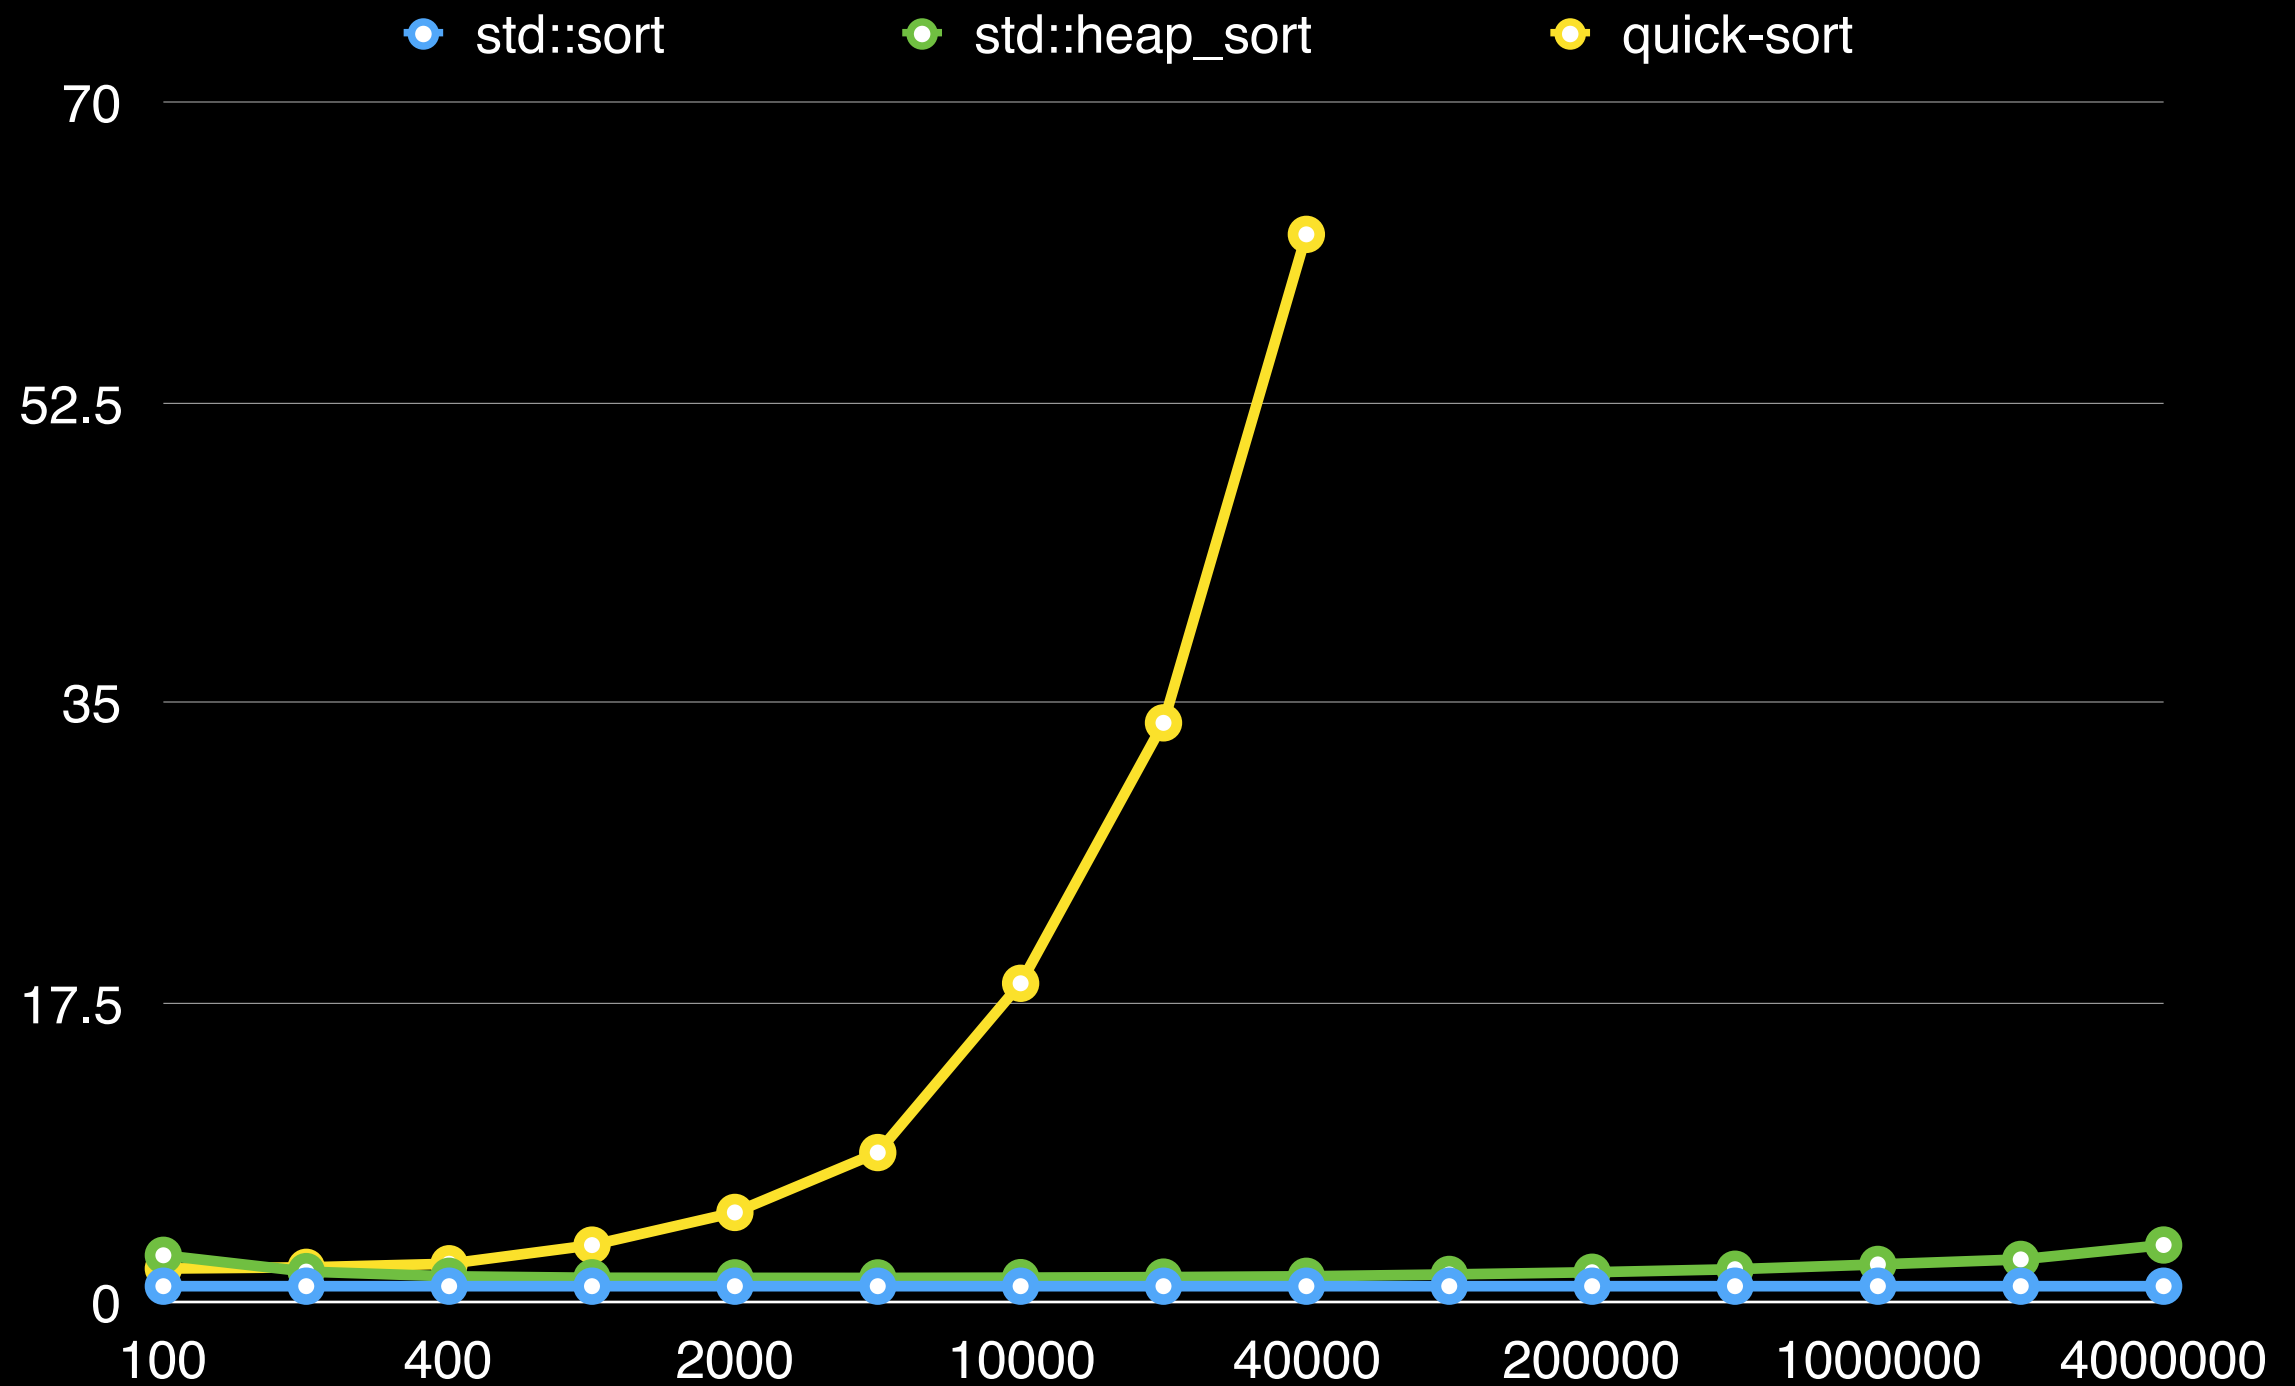
sort()

sort()

Simple Implementation

```
void sort(int* begin, int* end) {  
    std::size_t len = std::distance(begin, end);  
    if (len <= 1) return;  
    int* pivot = end - 1;  
    int* mid = std::partition(begin, pivot,  
                             [=](int value){ return value < *pivot; });  
    swap(*mid, *pivot);  
    sort(begin, mid);  
    sort(mid + 1, end);  
}
```

Results



Use a Pivot

- popular choice: median of 3 (or more)
- works well with sorted or reverse sorted inputs
- adds extra costs to determine the median
- ... using last is bad, though

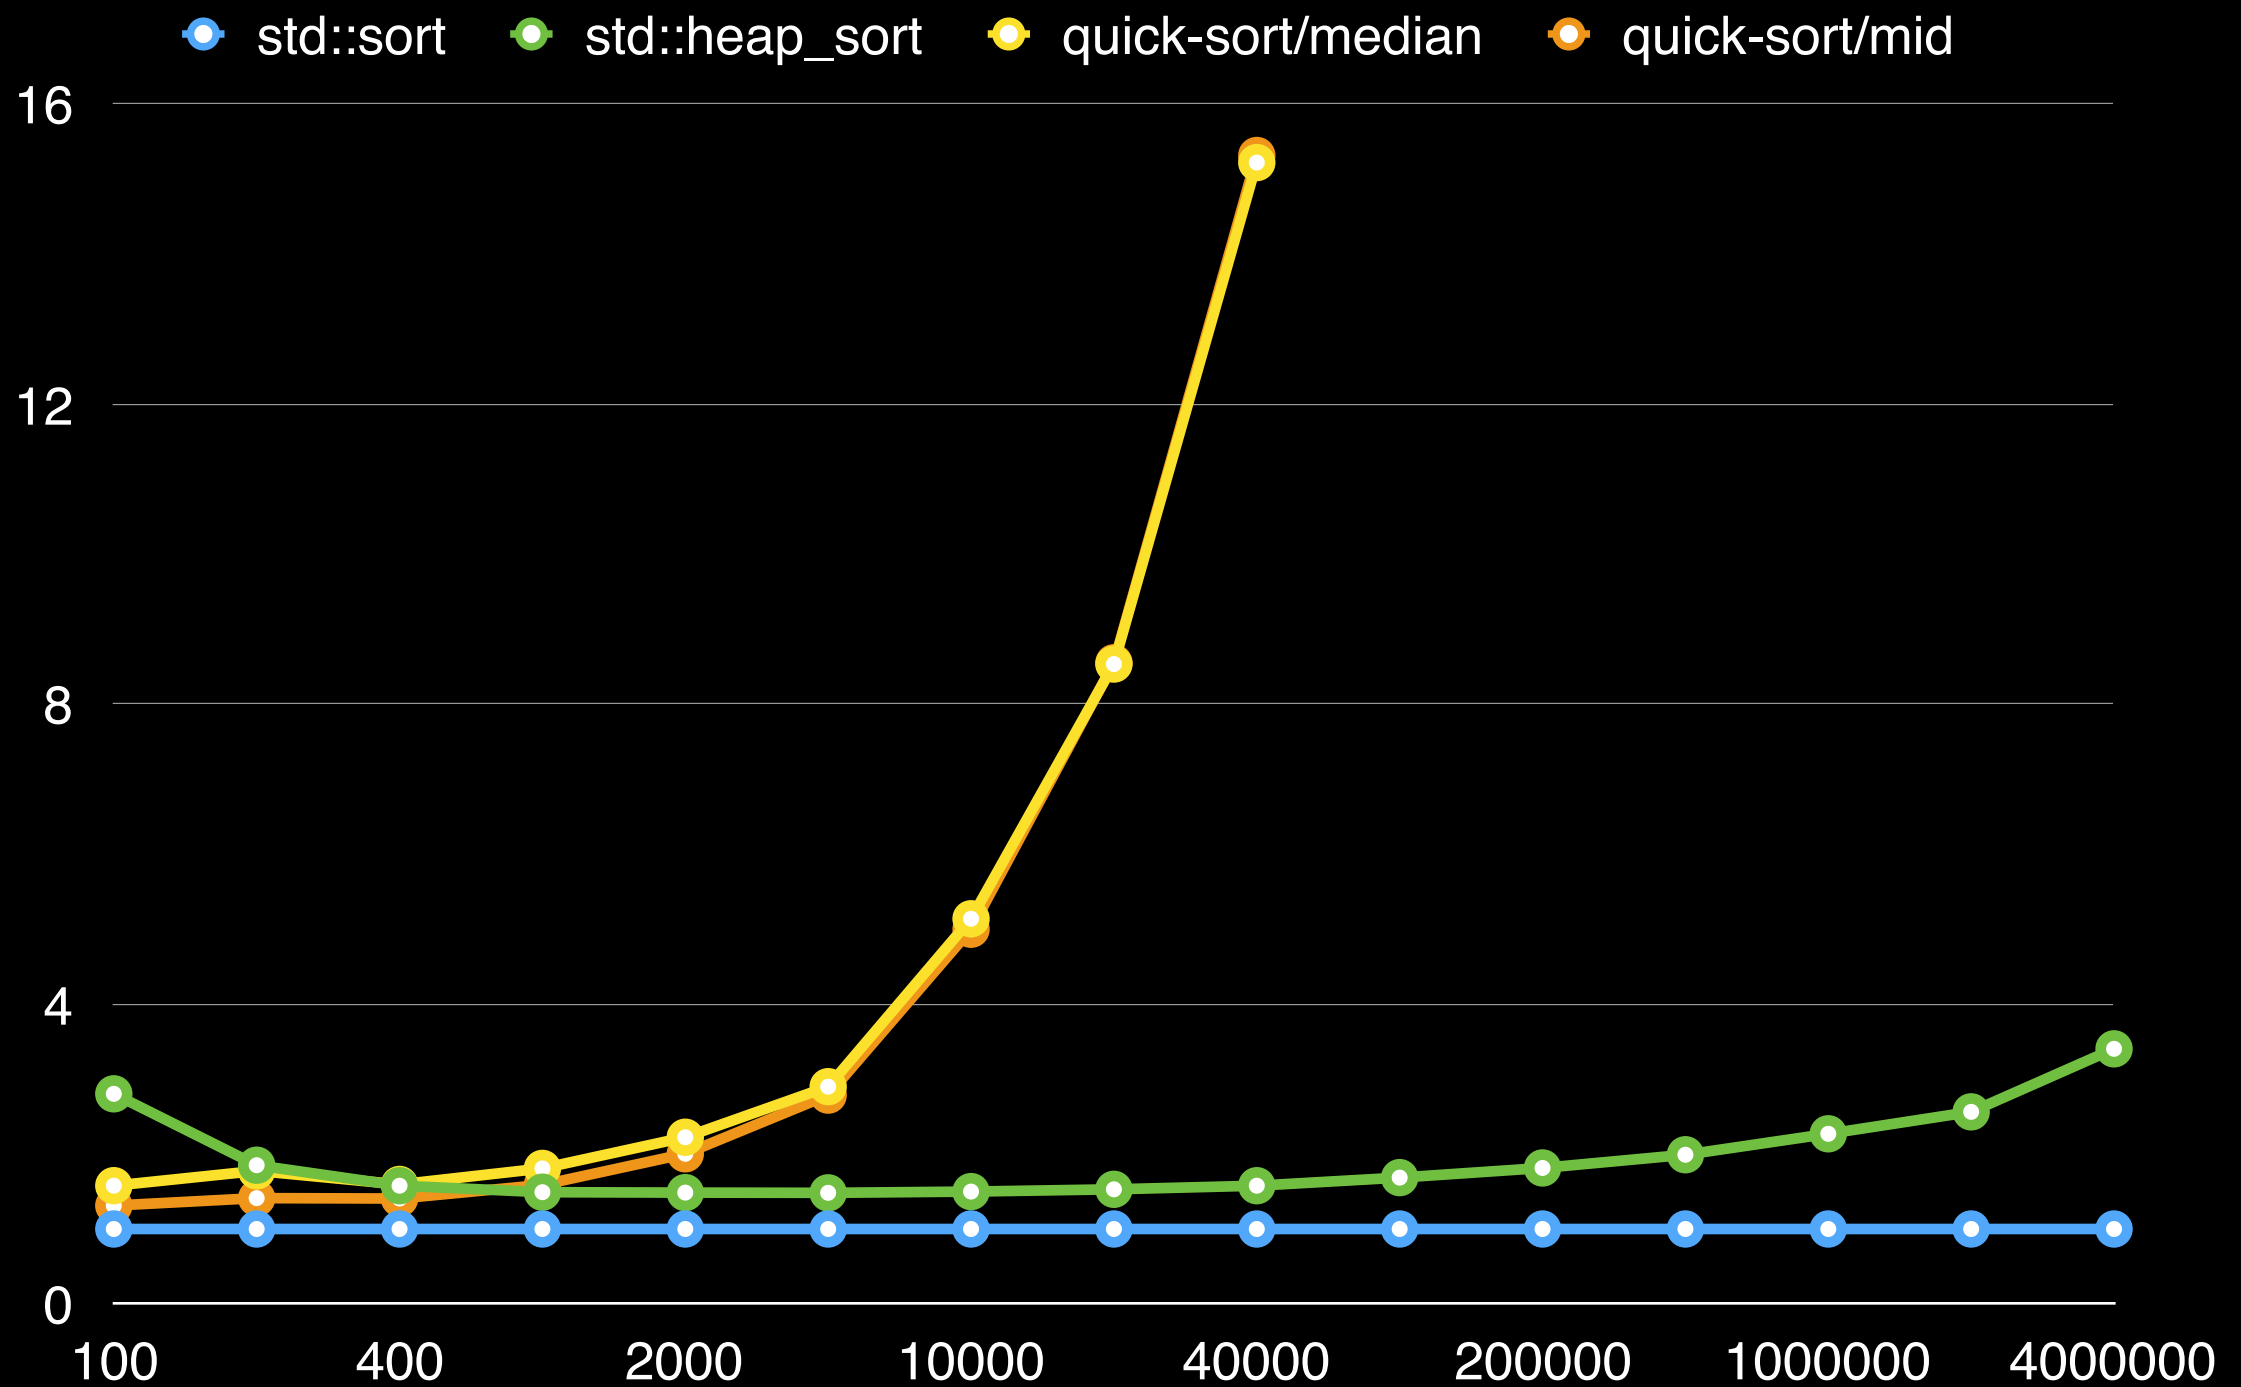
Pivot Implementation

```
void sort(int* begin, int* end) {  
    std::size_t len = std::distance(begin, end);  
    if (len <= 1) return;  
    int* pivot = end - 1, * mid = begin + len / 2;  
    small::sort(begin, pivot, mid);  
    mid = std::partition(begin, pivot,  
                        [=](int arg){ return arg < *pivot; });  
    swap(*mid, *pivot);  
    sort(begin, mid);  
    sort(mid + 1, end);  
}
```


Sorting 3 Elements

```
template <typename I>
static inline void small::sort(I a, I b, I c) {
    if (*b < *a) {
        if (*c < *b) std::iter_swap(a, c);
        else if (*c < *a) { auto t(*a); *a = *b; *b = *c; *c = t; }
        else std::iter_swap(a, b);
    } else {
        if (*c < *a) { auto t(*c); *c = *b; *b = *a; *a = t; }
        else if (*c < *b) std::iter_swap(b, c);
    }
}
```

Results



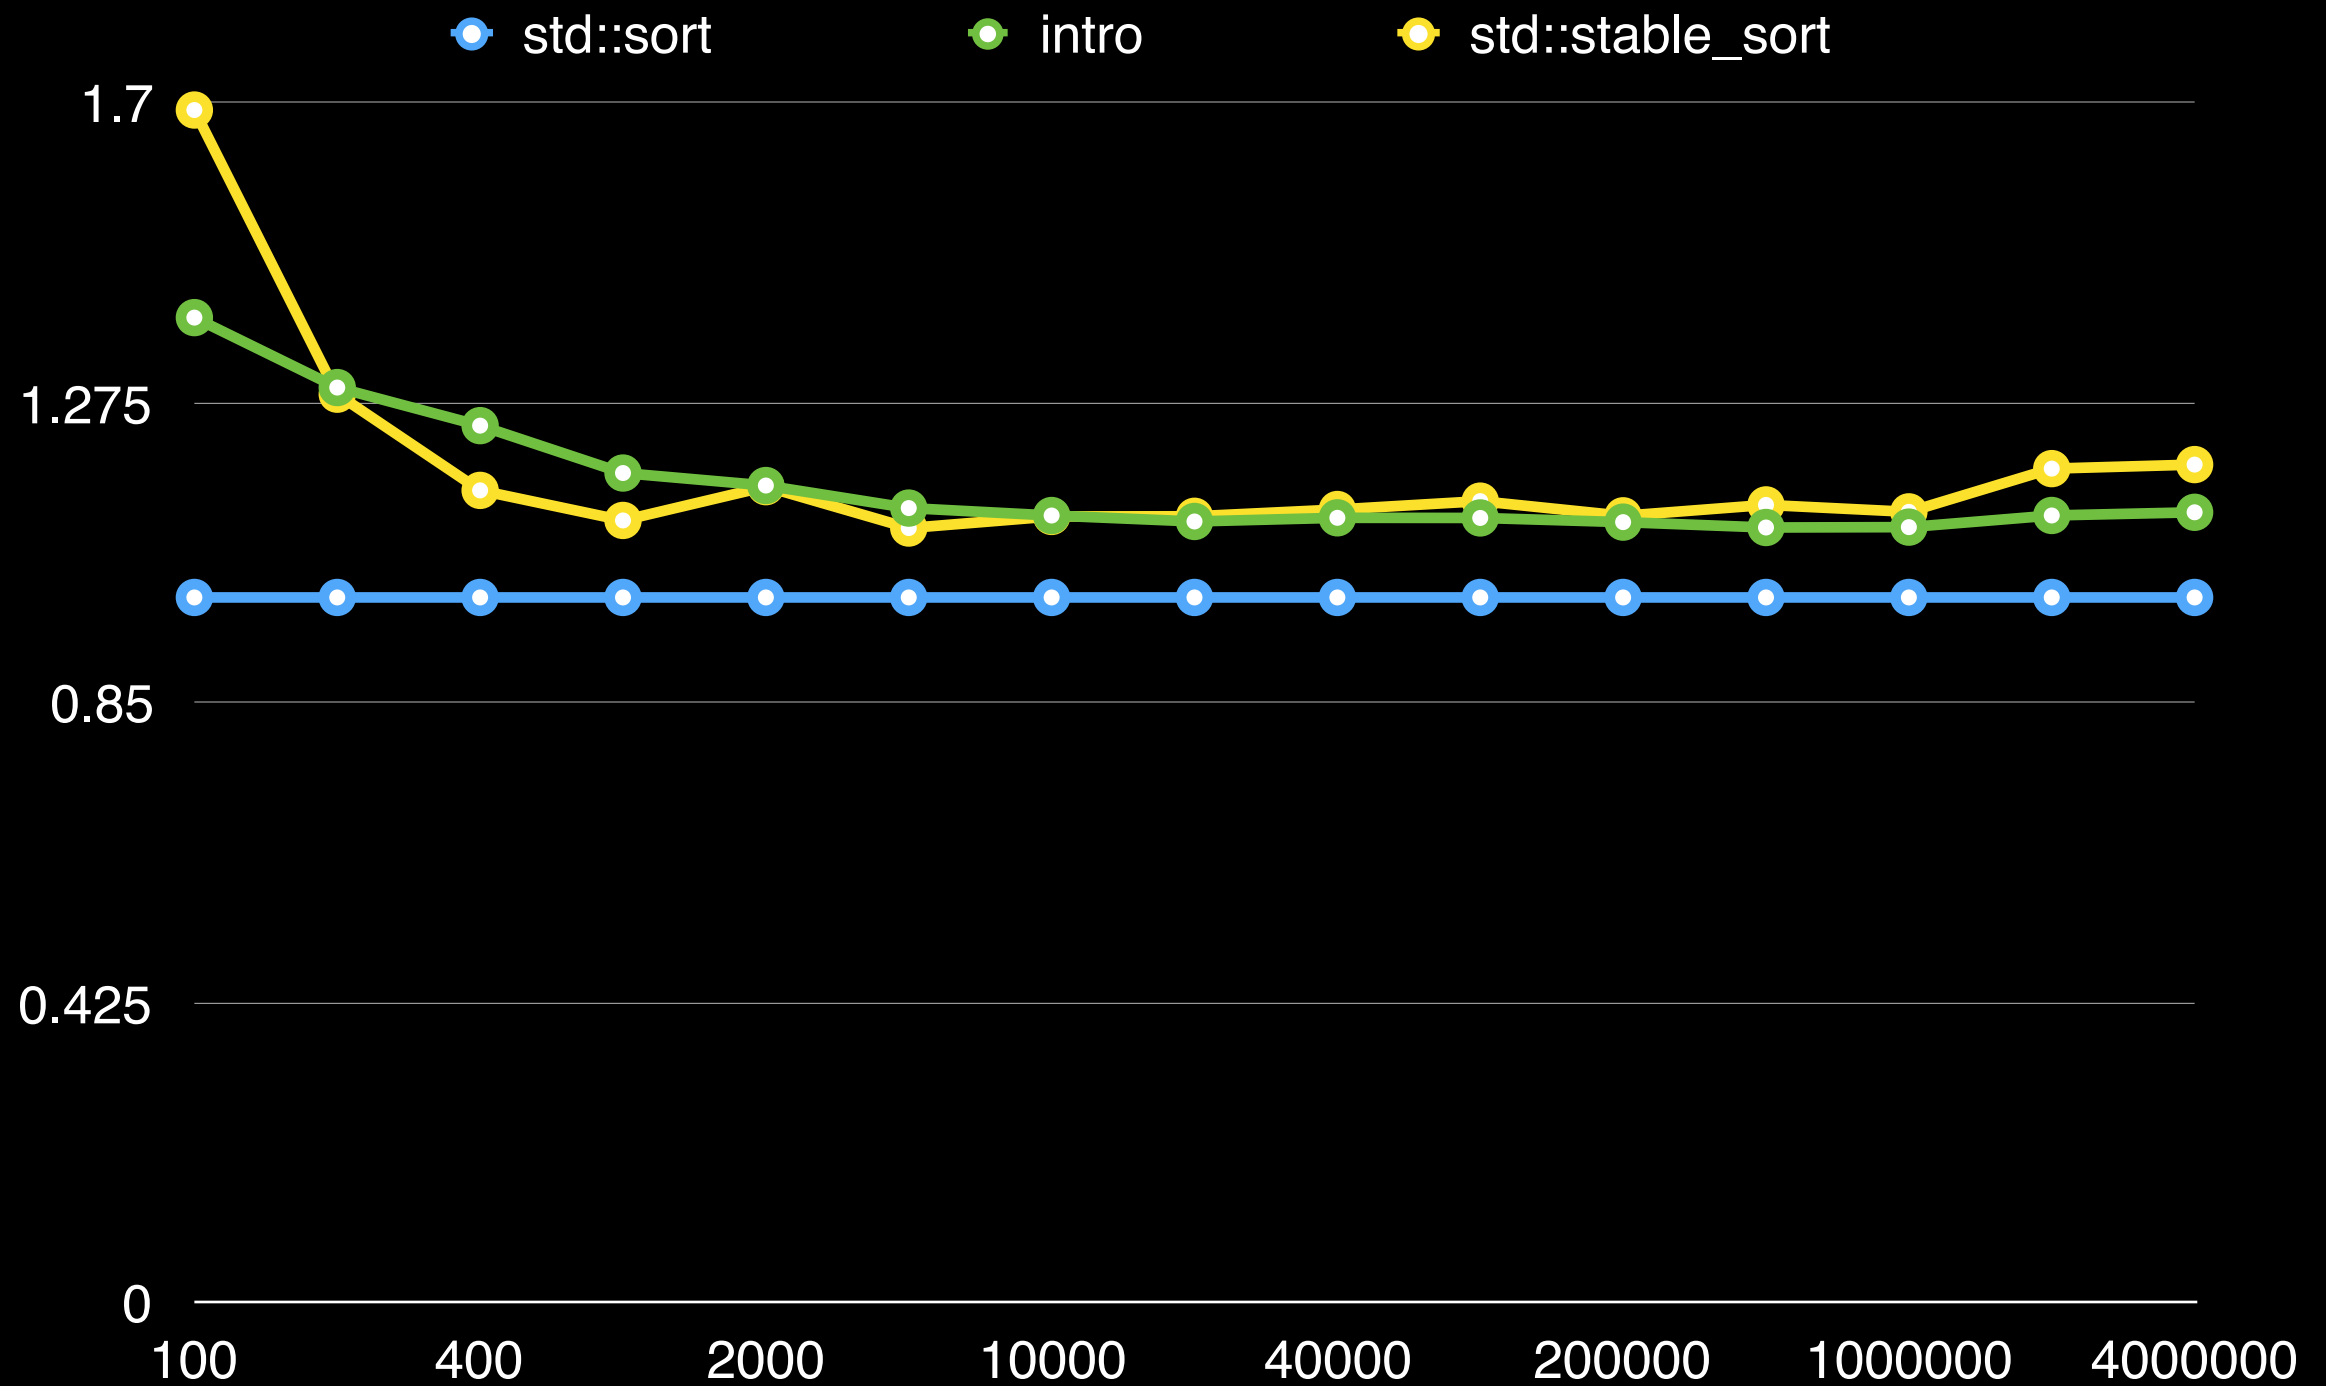
Intro Sort

- monitor depth of recursion:
 - no more than $2 \log n$ recursive calls
- too deep \Rightarrow fallback to a different algorithm
 - in-place required: heap sort
 - enough spare memory: merge sort

Intro Sort

```
void sort(int* begin, int* end, int depth, int max) {  
    if (++depth < max) {  
        ...  
        sort(begin, mid, depth, max);  
        sort(mid + 1, end, depth, max);  
    }  
    else { std::stable_sort(begin, end); }  
}  
  
void sort(int* begin, int* end) {  
    int s(end-begin), m(0); while (s >>= 1) { ++m; }  
    sort(begin, end, 0, 2 * m);  
}
```

Results




Insertion Sort

- quick sort isn't really effective on small ranges
- insertion sort is more effective in that case
- two potential approaches:
 - stop when ranges too small and post-process
 - sort small ranges with insertion sort

Local Insertion Sort

```
if (len <= Size) {  
    if (begin == end) return;  
    for (int* it = begin; ++it != end; )  
        if (*it < *(it - 1)) {  
            int* c = it, t = *c;  
            *c = *(c - 1);  
            while (begin != --c && t < *(c - 1))  
                *c = *(c - 1);  
            *c = t;  
        }  
    return;  
}
```

Results



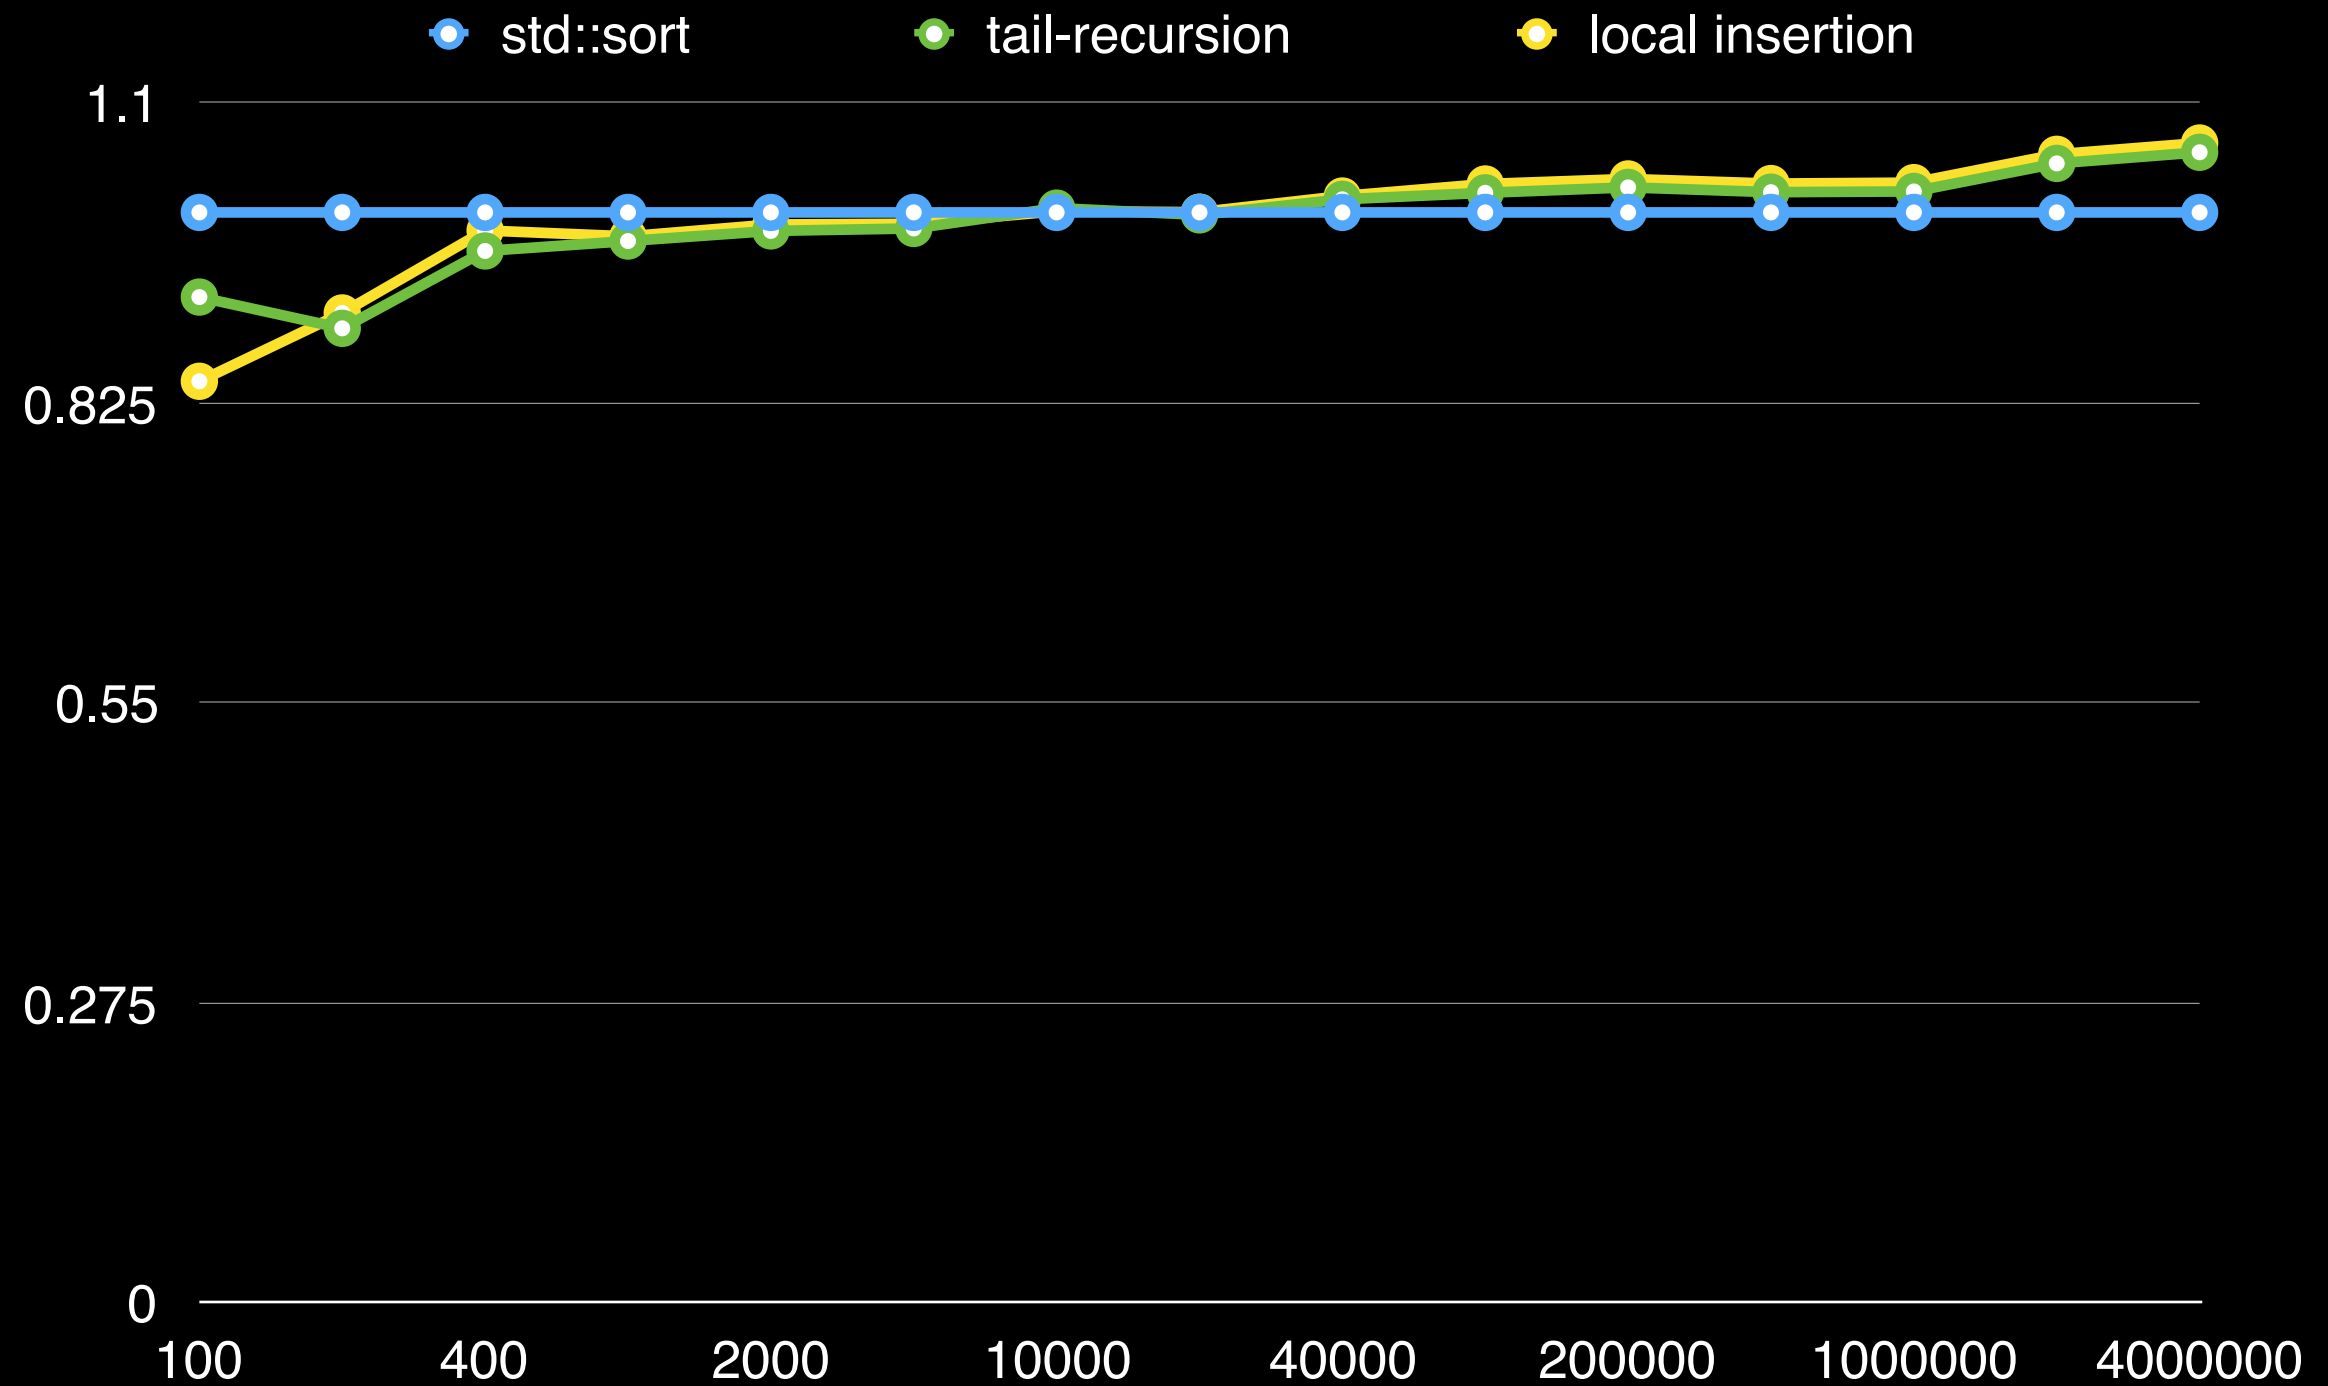
“Tail Recursion”

- using more stack is slower
- instead of call use a loop for the bigger part
- equivalent to tail call optimisation
 - however, C++ compilers are not too strong

Tail Recursion

```
while (20 < (len = std::distance(begin, end))) {  
    if (++depth < max) {  
        ...  
        if (mid - begin < end - mid) {  
            sort(begin, mid, depth, max);  
            begin = mid + 1; }  
        else {  
            sort(mid + 1, end, depth, max);  
            end = mid; }  
    }  
}
```

Results




Partition

- so far using `std::partition()`:
`mid = std::partition(begin, end, predicate);`
- really the core part of the algorithm


Partition

start




Partition


start




pivot




Partition




Partition




Partition




Partition




Partition



Partition



Partition



Partition

```
...  
int* pbegin = begin, * pend = end;  
while (true) {  
    while (pbegin != pend && *pbegin < *pivot)  
        ++pbegin;  
    while (pbegin != pend && !(*--pend < *pivot));  
    if (pbegin == pend) break;  
    std::iter_swap(pbegin, pend);  
    ++pbegin;  
}  
mid = pbegin;  
...
```

Results



Sentinel Partition

- inner loop makes two checks per iteration:
 - determine if the end is reached
 - determine if the predicate applies
- can the check for the end be avoided?
 - use an object matching the predicate to stop

Sentinel on the Right

```
...
int* pbegin = begin, * pend = end;
while (true) {
    while (pbegin != pend && *pbegin < *pivot)
        ++pbegin;
    while (pbegin != pend && !(*--pend < *pivot));
    if (pbegin == pend) break;
    std::iter_swap(pbegin, pend);
    ++pbegin;
}
mid = pbegin;
...
```

Sentinel on the Right

```
...  
int* pbegin = begin, * pend = end;  
while (true) {  
    while (*pbegin < *pivot)  
        ++pbegin;  
    while (pbegin != pend && !(*--pend < *pivot));  
    if (pbegin == pend) break;  
    std::iter_swap(pbegin, pend);  
    ++pbegin;  
}  
mid = pbegin;  
...
```

Initial Sentinel Partition


make sure there is a sentinel on the left side, too:

```
int* pbegin = begin, * pend = end;
if (pbegin != pend && !(*pbegin < p)) {
    while (pbegin != pend && !(*--pend < p)) {
    }
    if (pbegin != pend) {
        std::iter_swap(pbegin, pend);
        ++pbegin;
    }
}
```

Sentinel Partition Core

```
if (pbegin != pend) {  
    while (true) {  
        while (*pbegin < p) ++pbegin;  
        while (!(*--pend < p));  
        if (pend <= pbegin) break;  
        std::iter_swap(pbegin, pend);  
        ++pbegin;  
    }  
}  
mid = pbegin;
```

Results



Small Ranges

- many of the final ranges are small
- not worth kicking off an insertion sort
- manually sort ranges with up to 4 elements
- bigger ranges still use normal insertion sort


Small Range Dispatch

```
switch (std::distance(b, e)) {  
  case 0:  
  case 1: return;  
  case 2: if (*--e < *b) swap(*b, *e); return;  
  case 3: small::sort(b, b + 1, b + 2); return;  
  case 4: small::sort(b, b + 1, b + 2, b + 3); return;  
}
```

Sorting Four Elements

```
void small::sort(l p0, l p1, l p2, l p3) {  
    sort(p0, p1, p2);  
    if (*p3 < *p2) {  
        if (*p3 < *p1) {  
            auto tmp = *p3; *p3 = *p2; *p2 = *p1;  
            if (tmp < *p0) { *p1 = *p0; *p0 = tmp; }  
            else *p1 = tmp;  
        }  
        else swap(p2, p3);  
    }  
}
```


Results



Partition Did Nothing

- the partition may not do anything
 - sequence is sorted already
 - notably: all elements are the same
- gamble time on testing if it is sorted
 - ... and fix a few local inconsistencies

Monitor Swaps


splitting the check may trigger small range optimisation

```
unsigned swaps = 0;  
... // partition incrementing swaps  
  
if (swaps == 0  
    && incomplete_insertion(begin, mid)  
    && incomplete_insertion(mid + 1, end)) {  
    return;  
}
```

Partial Insertion Sort

```
switch (std::distance(b, e)) { ... }  
for (int* it = b, cout = 0; ++it != e; )  
    if (*it < *(it - 1)) {  
        if (++count == limit) return false;  
        int* c = it, t = *c; *c = *(c - 1);  
        while (b != --c && t < *(c - 1)) {  
            *c = *(c - 1);  
            if (++count == limit) { *(c - 1) = t; return false; }  
        }  
        *c = t;  
    }  
return true;
```

Results



Putting It All Together

```
bool incomplete_insertion(int* begin, int* end) {
    switch (std::distance(begin, end)) {
    case 0:
    case 1:
        return true;
    case 2:
        if (*--end < *begin) std::swap(*begin, *end);
        return true;
    case 3:
        small::sort(begin, begin + 1, begin + 2);
        return true;
    case 4:
        small::sort(begin, begin + 1, begin + 2, begin + 3);
        return true;
    case 5:
        small::sort(begin, begin + 1, begin + 2, begin + 3, begin
+ 4);
        return true;
    }
    constexpr int limit = 8;
    int count(0);
    for (int* it = begin; ++it != end; )
        if (*it < *(it - 1)) {
            if (++count == limit) return false;
            int* c = it, t = *c;
            *c = *(c - 1);
            while (begin != --c && t < *(c - 1)) {
                *c = *(c - 1);
                if (++count == limit) {
                    *(c - 1) = t;
                    return false;
                }
            }
            *c = t;
        }
    return true;
}

void sort(int* begin, int* end, int depth, int max) {
    std::size_t len;
    while (20 < (len = std::distance(begin, end))) {
        if (++depth < max) {
            int* pivot = end - 1, * mid = begin + len / 2;
```

```
            std::swap(*mid, *pivot);

            unsigned swaps = 0;
            int p = *pivot;
            auto pred = [](int arg){ return arg < p; };
            int* pbegin = begin, * pend = end;
            if (pbegin != pend && !pred(*pbegin)) {
                while (pbegin != pend && !pred(*--pend)) {
                }
                if (pbegin != pend) {
                    std::iter_swap(pbegin, pend);
                    ++swaps;
                    ++pbegin;
                }
            }
            if (pbegin != pend) {
                while (true) {
                    while (pred(*pbegin)) {
                        ++pbegin;
                    }
                    while (!pred(*--pend)) {
                    }
                    if (pend <= pbegin) {
                        break;
                    }
                    std::iter_swap(pbegin, pend);
                    ++pbegin;
                    ++swaps;
                }
            }
            mid = pbegin;
            std::swap(*mid, *pivot);

            if (swaps == 0
                && incomplete_insertion(begin, mid)
                && incomplete_insertion(mid + 1, end)) {
                return;
            }
            if (mid - begin < end - mid) {
                sort(begin, mid, depth, max);
                begin = mid + 1;
```

```
            }
            else {
                sort(mid + 1, end, depth, max);
                end = mid;
            }
        }
        else {
            std::stable_sort(begin, end);
            return;
        }
    }
    switch (std::distance(begin, end)) {
    case 0:
    case 1:
        return;
    case 2:
        if (*--end < *begin) std::swap(*begin, *end);
        return;
    case 3:
        small::sort(begin, begin + 1, begin + 2);
        return;
    case 4:
        small::sort(begin, begin + 1, begin + 2, begin + 3);
        return;
    }
    for (int* it = begin; ++it != end; )
        if (*it < *(it - 1)) {
            int* c = it, t = *c;
            *c = *(c - 1);
            while (begin != --c && t < *(c - 1))
                *c = *(c - 1);
            *c = t;
        }
    return;
}

void sort(int* begin, int* end) {
    std::size_t size(std::distance(begin, end)), max(0);
    while (size >>= 1) { ++max; }
    sort(begin, end, 0, 2 * max);
}
```

Putting It All Together

```
bool incomplete_insertion(int* begin, int* end) {
    switch (std::distance(begin, end)) {
    case 0:
    case 1:
        return true;
    case 2:
        if (--end < *begin) std::swap(*begin, *end);
        return true;
    case 3:
        small::sort(begin, begin + 1, begin + 2);
        return true;
    case 4:
        small::sort(begin, begin + 1, begin + 2, begin + 3);
        return true;
    case 5:
        small::sort(begin, begin + 1, begin + 2, begin + 3, begin
+ 4);
        return true;
    }
    constexpr int limit = 8;
    int count(0);
    for (int* it = begin; ++it != end; )
        if (*it < *(it - 1)) {
            if (++count == limit) return false;
            int* c = it, t = *c;
            *c = *(c - 1);
            while (begin != --c && t < *(c - 1)) {
                *c = *(c - 1);
                if (++count == limit) {
                    *(c - 1) = t;
                    return false;
                }
            }
            *c = t;
        }
    return true;
}

void sort(int* begin, int* end, int depth, int max) {
    std::size_t len;
    while (20 < (len = std::distance(begin, end))) {
        if (++depth < max) {
            int* pivot = end - 1, * mid = begin + len / 2;
```

```
            std::swap(*mid, *pivot);

            unsigned swaps = 0;
            int p = *pivot;
            auto pred = [=](int arg){ return arg < p; };
            int* pbegin = begin, * pend = end;
            if (pbegin != pend && !pred(*pbegin)) {
                while (pbegin != pend && !pred(*--pend)) {
                }
                if (pbegin != pend) {
                    std::iter_swap(pbegin, pend);
                    ++swaps;
                    ++pbegin;
                }
            }
            if (pbegin != pend) {
                while (true) {
                    while (pred(*pbegin)) {
                        ++pbegin;
                    }
                    while (!pred(*--pend)) {
                    }
                    if (pend <= pbegin) {
                        break;
                    }
                    std::iter_swap(pbegin, pend);
                    ++pbegin;
                    ++swaps;
                }
            }
            mid = pbegin;
            std::swap(*mid, *pivot);

            if (swaps == 0
                && incomplete_insertion(begin, mid)
                && incomplete_insertion(mid + 1, end)) {
                return;
            }
            if (mid - begin < end - mid) {
                sort(begin, mid, depth, max);
                begin = mid + 1;
```

```
            }
            else
                sort(mid + 1, end, depth, max);
            end = mid;
        }
    }
    else {
        std::stable_sort(begin, end);
        return;
    }
}

switch (std::distance(begin, end)) {
case 0:
case 1:
    return;
case 2:
    if (--end < *begin) std::swap(*begin, *end);
    return;
case 3:
    small::sort(begin, begin + 1, begin + 2);
    return;
case 4:
    small::sort(begin, begin + 1, begin + 2, begin + 3);
    return;
}
for (int* it = begin; ++it != end; )
    if (*it < *(it - 1)) {
        int* c = it, t = *c;
        *c = *(c - 1);
        while (begin != --c && t < *(c - 1))
            *c = *(c - 1);
        *c = t;
    }
return;
}

void sort(int* begin, int* end) {
    std::size_t size(std::distance(begin, end)), max(0);
    while (size >>= 1) { ++max; }
    sort(begin, end, 0, 2 * max);
}
```

Simple Implementation

```
void sort(int* begin, int* end) {  
    std::size_t len = std::distance(begin, end);  
    if (len <= 1) return;  
    int* pivot = end - 1;  
    int* mid = std::partition(begin, pivot,  
                             [=](int value){ return value < *pivot; });  
    swap(*mid, *pivot);  
    sort(begin, mid);  
    sort(mid + 1, end);  
}
```


More Generic

```
void sort(int* begin, int* end) {  
    std::size_t len = std::distance(begin, end);  
    if (len <= 1) return;  
    int* pivot = end - 1;  
    int* mid = std::partition(begin, pivot,  
                             [=](int value){ return value < *pivot; });  
    swap(*mid, *pivot);  
    sort(begin, mid);  
    sort(mid + 1, end);  
}
```

More Generic

```
template <typename T>
void sort(T * begin, T * end) {
    std::size_t len = std::distance(begin, end);
    if (len <= 1) return;
    T * pivot = end - 1;
    T * mid = std::partition(begin, pivot,
                             [=](T value){ return value < *pivot; });
    swap(*mid, *pivot);
    sort(begin, mid);
    sort(mid + 1, end);
}
```

More Generic

```
template <typename Iter>
void sort(Iter begin, Iter end) {
    std::size_t len = std::distance(begin, end);
    if (len <= 1) return;
    Iter pivot = end - 1;
    Iter mid = std::partition(begin, pivot,
        [=](auto&& value){ return value < *pivot; });
    swap(*mid, *pivot);
    sort(begin, mid);
    sort(mid + 1, end);
}
```

More Generic

```
template <typename Iter>
void sort(Iter begin, Iter end) {
    std::size_t len = std::distance(begin, end);
    if (len <= 1) return;
    Iter pivot = end - 1;
    Iter mid = std::partition(begin, pivot,
        [=](auto&& value){ return value < *pivot; });
    swap(*mid, *pivot);
    sort(begin, mid);
    sort(mid + 1, end);
}
```

More Generic

```
template <typename Iter, typename Comp>
void sort(Iter begin, Iter end, Comp c) {
    std::size_t len = std::distance(begin, end);
    if (len <= 1) return;
    Iter pivot = end - 1;
    Iter mid = std::partition(begin, pivot,
        [=](auto&& value){ return c(value, *pivot); });
    swap(*mid, *pivot);
    sort(begin, mid, c);
    sort(mid + 1, end, c);
}
```

Quicksort Killer

- custom comparators allow hostile orders:
 - make every choice of pivot bad
 - adaptive order based on elements compared
 - have the pivot always be rather small

Killer Comparison

```
bool operator()(int x, int y) const {  
    if (r->v[x] == max && r->v[y] == max) {  
        r->v[x == r->cand? x: y] = ++r->count;  
        r->cand = x == r->cand? y: x;  
    }  
    return r->v[x] < r->v[y];  
}
```

Killer Comparison

```
bool operator()(int x, int y) const {  
    if (r->v[x] == max && r->v[y] == max) {  
        r->v[x == r->cand? x: y] = ++r->count;  
        r->cand = x == r->cand? y: x;  
    }  
    return r->v[x] < r->v[y];  
}
```


Killer Comparison

```
bool operator()(int x, int y) const {  
    if (r->v[x] == max && r->v[y] == max) {  
        r->v[x == r->cand? x: y] = ++r->count;  
        r->cand = x == r->cand? y: x;  
    }  
    return r->v[x] < r->v[y];  
}
```

Summary

- quick sort is fast
 - when pulling a lot of tricks
 - for the expected cases
 - it is a hybrid using multiple algorithms
- => real-world algorithms need a bit of tuning

Questions

